

# 演示：Service创建通路

展示通过kubectl创建service，到ServiceController分配ClusterIP，返回给apiserver，最后通知proxy的流程

信息通路：kubectl->apiserver->ServiceController->apiserver->proxy

## 流程

首先通过kubectl创建一个pod

```
→ build git:(kubelet) X ./kubectl apply -f pod-example.yaml
ok
→ build git:(kubelet) X cat pod-example.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod
  labels:
    app: nginx
spec:
  restartPolicy: Always
  containers:
    - name: viewer
      image: dplsming/nginx-fileserver:1.0
      ports:
        - containerPort: 80
          hostPort: 8888
      volumeMounts:
        - name: volume
          mountPath: /usr/share/nginx/html/files
    - name: downloader
      image: dplsming/aria2ng-downloader:1.0
      ports:
        - name: nginx
          containerPort: 6800
          hostPort: 6800
        - name: nginx
          containerPort: 6880
          hostPort: 6880
      volumeMounts:
        - name: volume
          mountPath: /data
  volumes:
    - name: volume
      hostPath:
        path: /pod
    - name: nfs-volume
      nfs:
        path: /exports
        server: 192.168.10.1
```

然后再创建一个Service

```
→ build git:(kubelet) X ./kubectl apply -f service.yaml
ok
→ build git:(kubelet) X cat service.yaml
apiVersion: v1
kind: Service
metadata:
  name: myService
  namespace: default
spec:
  selector:
    app: nginx
  type: ClusterIP
  ports:
    - name: myPort
      protocol: TCP
      port: 8080 # 对外暴露的端口
      targetPort: nginx # 转发的端口的名字, pod对应的端口名字
```

- apiserver收到http请求, 将service信息告知ServiceController, 并等待ServiceController分配ClusterIP
- ServiceController收到service创建的通知, 根据策略, 分配ClusterIP, 然后筛选满足条件的pod, 并生成对应的endpoints信息, 将带有分配好的ClusterIP的service信息以及筛选出的endpoints信息回传给apiserver
- apiserver收到带有分配好的ClusterIP的service信息以及筛选出的endpoints信息, 将其存入etcd, 并发布service创建的通知
- proxy监听到service创建的通知, 调用ipvs指令创建service及其对应的endpoints

```
swung@swung-Lenovo-XiaoXinAir-15ITL-2021:~/桌面/minik8s/build$ ./kubectl apply -f ../kubectl/run/pod-example.yaml
ok
swung@swung-Lenovo-XiaoXinAir-15ITL-2021:~/桌面/minik8s/build$ ./kubectl apply -f ../apiobjects/examples/pod/service.yaml
apply service: {v1 {myService default} {ClusterIP [{myPort TCP 8080 nginx}] map[app:nginx]} {}}
ok
swung@swung-Lenovo-XiaoXinAir-15ITL-2021:~/桌面/minik8s/build$
```

上图为使用kubectl分别创建一个pod和service

```
[GIN-debug] Listening and serving HTTP on :8080
receive pod name: pod2 namespace: default uid: 7837cc4d-8753-449e-a6fc-2a0f36ba6596[GIN] 2024/05/13 - 14:59:37 | 200 | 215.648173ms | 127.0.0.1 | POST "/api/pod"
receive service name: myService namespace: default uid: f05a22cb-6548-4f0d-9549-5d739e03ca25[GIN] 2024/05/13 - 15:00:18 | 200 | 58.079774ms | 127.0.0.1 | POST "/api/service/apply"
service create: {"apiVersion":"v1","metadata":{"name":"myService","namespace":"default","uid":"f05a22cb-6548-4f0d-9549-5d739e03ca25"},"spec":{"type":"ClusterIP","ports":[{"name":"myPort","protocol":"TCP","port":8080,"targetPort":{"app":"nginx"}},{"selector":{"app":"nginx"},"status":{"phase":"CREATED","clusterIP":"10.10.0.1"}}]}
[GIN] 2024/05/13 - 15:00:18 | 200 | 142.064279ms | 127.0.0.1 | POST "/api/service"
INFO[0070] pod value: {"ApiVersion":"v1","Kind":"Pod","Name":"pod2","Namespace":"default","UID":"7837cc4d-8753-449e-a6fc-2a0f36ba6596","Labels":{"app":"nginx"},"CreationTimestamp":"2024-05-13T14:59:37.435182136+08:00","DeletionTimestamp":"0001-01-01T00:00:00Z","Spec":{"Containers":[{"Name":"viewer","Image":"dplsming/nginx-fileserver:1.0","Ports":[{"Name":"","ContainerPort":80,"HostPort":8888}], "VolumeMounts":[{"Name":"volume","MountPath":"/usr/share/nginx/html/files"}], "Labels":null, "Status":null}, {"Name":"downloader","Image":"dplsming/aria2ng-downloader:1.0","Ports":[{"Name":"nginx","ContainerPort":6880,"HostPort":6880}, {"Name":"nginx","ContainerPort":6880,"HostPort":6880}], "VolumeMounts":[{"Name":"volume","MountPath":"/data"}], "Labels":null, "Status":null}, {"Name":"volume","EmptyDir":null, "HostPath":{"Path":"/pod"},"NFS":null, "PersistentVolumeClaim":null}], "NodeSelector":null}, {"Status":{"PodPhase":"","HostIP":"","PodIP":""}}}
[GIN] 2024/05/13 - 15:00:18 | 200 | 5.143269ms | 127.0.0.1 | GET "/api/get/allpods"
endpoint create: {"serviceName":"myService","metadata":{"name":"myService-pod2","namespace":"default"},"spec":{"svcIP":"10.10.0.1","svcPort":8080,"dstIP":"","dstPort":6880}}
[GIN] 2024/05/13 - 15:00:18 | 200 | 155.748667ms | 127.0.0.1 | POST "/api/endpoint"
```

上图显示apiserver首先从kubectl收到“/api/service/apply”的消息, 然后将消息告知ServiceController后, ServiceController分配ClusterIP之后, 通过“/api/service”再将带有分配好的ClusterIP的service信息回传给apiserver, 同时, ServiceController还要筛选满足条件的pod, ServiceController首先通过“/api/get/allpods”从apiserver拿到所有的pod信息, 然后进行筛选, 并生成对应的endpoints信息,

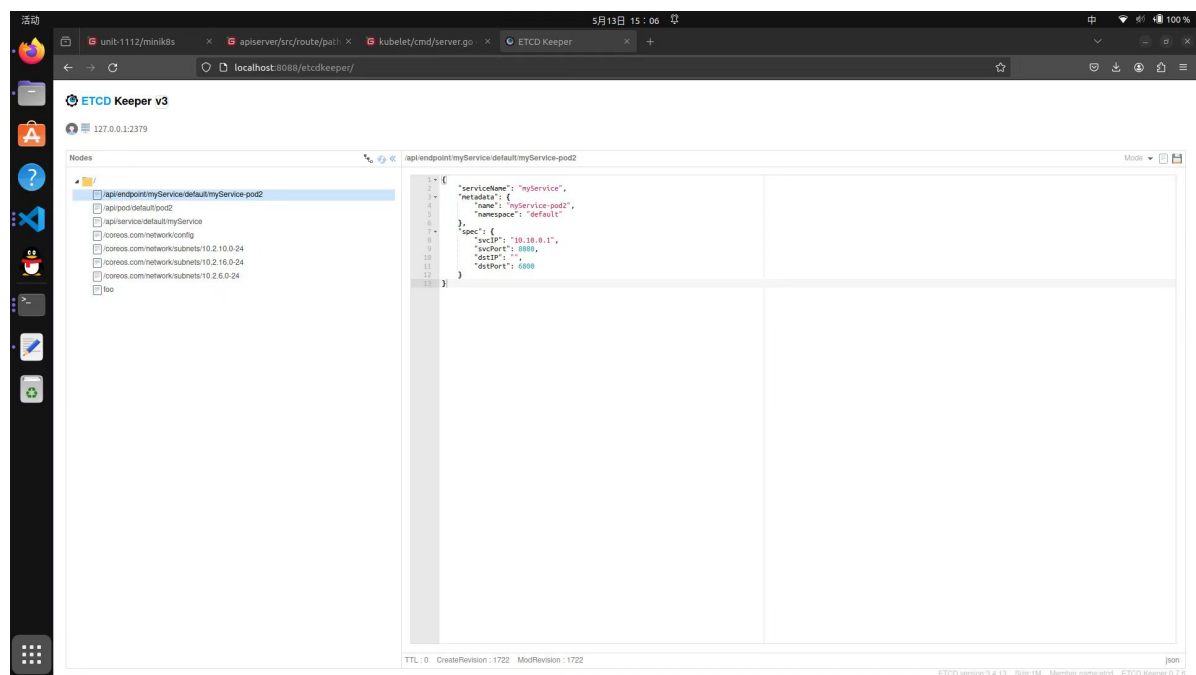
最后通过“api/endpoint”将筛选出的endpoints信息回传给apiserver

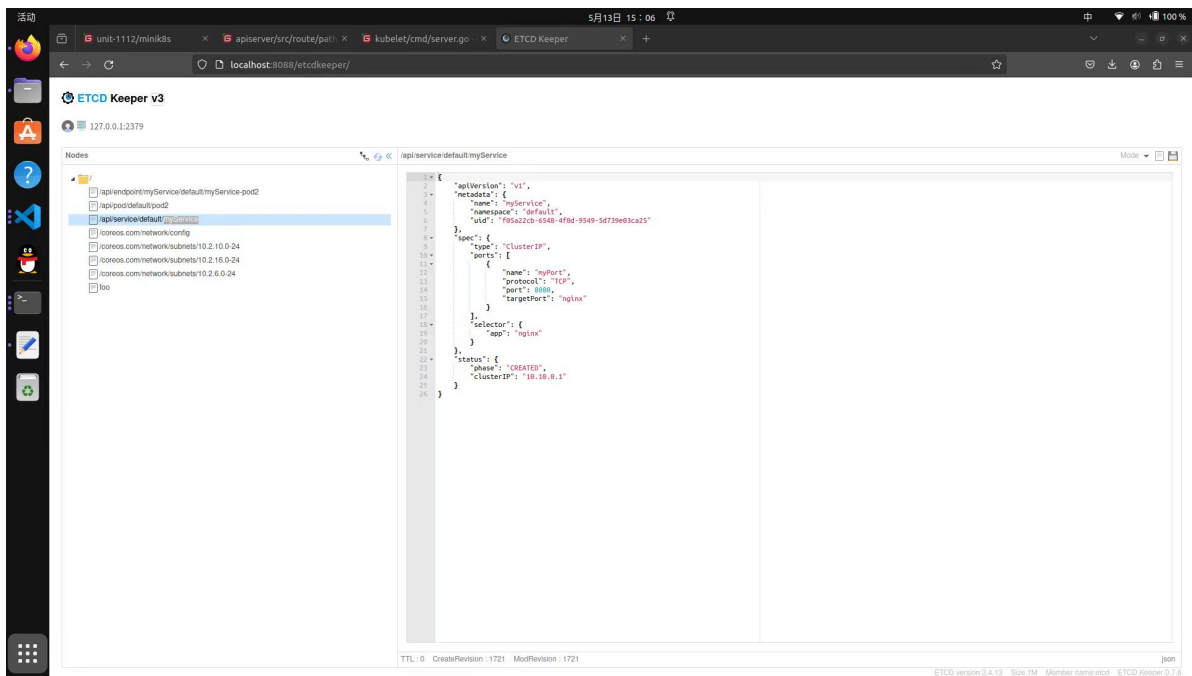
```
swung@swung-Lenovo-XiaoXinAir-15ITL-2021:~/桌面/minik8s/build$ ./controllermanager
msg-payload:{"ActionType":0,"Object":{"apiVersion":"v1","metadata":{"name":"myService","namespace":"default","uid":"f05a22cb-6548-4f0d-9549-5d739e03ca25"},"spec":{"type":"ClusterIP","ports":[{"name":"myPort","protocol":"TCP","port":8080,"targetPort":"nginx"}],"selector":{"app":"nginx"},"status":{"phase":"CREATING","clusterIP":"","ip":"10.10.0.1"}}, msg-channel:service-cmdHandleServiceApply
HandleServiceApply myService
8{200 OK 200 HTTP/1.1 1 1 map[Content-Length:[2] Content-Type:[text/plain; charset=utf-8] Date:[Mon, 13 May 2024 07:00:18 GMT]] 0xc000130040 2 [] false false map[
] 0xc0000b2ea0 <nil>}
INFO[0058] GetUnmarshal[{"ApiVersion":"v1","Kind":"Pod","Name":"pod2","Namespace":"default","UID":"7837cc4d-8753-449e-a6fc-2a0f36ba6596","Labels":{"app":"nginx"},"
"CreationTimestamp":"2024-05-13T14:59:37.435182136+08:00","DeletionTimestamp":"0001-01-01T00:00:00Z","Spec":{"Containers":[{"Name":"viewer","Image":"dplsming/nginx
x-fileserver:1.0","Ports":[{"Name":"","ContainerPort":80,"HostPort":8888}], "VolumeMounts":[{"Name":"volume","MountPath":"/usr/share/nginx/html/files}], "Labels":n
ull,"Status":null}, {"Name":"downloader","Image":"dplsming/aria2ng-downloader:1.0","Ports":[{"Name":"nginx","ContainerPort":6800,"HostPort":6800}, {"Name":"nginx","
ContainerPort":6880,"HostPort":6880}], "VolumeMounts":[{"Name":"volume","MountPath":"/data"}], "Labels":null,"Status":null}, {"Name":"volume","EmptyDir":
null,"HostPath":{"Path":"/pod"},"NFS":null,"PersistentVolumeClaim":null}], "NodeSelector":null,"Status":{"PodPhase":"","HostIP":"","PodIP":""}}]
8{200 OK 200 HTTP/1.1 1 1 map[Content-Length:[2] Content-Type:[text/plain; charset=utf-8] Date:[Mon, 13 May 2024 07:00:18 GMT]] 0xc000130200 2 [] false false map[
] 0xc0001f2fc0 <nil>}
INFO[0058] [svc controller] Create endpoints.srcIP:10.10.0.1:8080, dstIP:6800 ;
INFO[0058] [svc controller] Create service. Cluster IP:10.10.0.1
```

上图显示ServiceController为新创建的service分配的ClusterIP以及筛选出的满足条件的pod对应创建的endpoint信息

```
swung@swung-Lenovo-XiaoXinAir-15ITL-2021:~/桌面/minik8s/build$ ./proxy
msg-payload:{"ActionType":0,"Object":{"apiVersion":"v1","metadata":{"name":"myService","namespace":"default","uid":"f05a22cb-6548-4f0d-9549-5d739e03ca25"},"spec":{"type":"ClusterIP","ports":[{"name":"myPort","protocol":"TCP","port":8080,"targetPort":"nginx"}],"selector":{"app":"nginx"},"status":{"phase":"CREATED","clusterIP":"10.10.0.1"}}, msg-channel:service10.10.0.1:8080NMsgerr operation not permitted
exit status 1
exit status 2
INFO[0055] [kubeproxy] Add service 10.10.0.1:8080
```

上图显示proxy创建对应的service





etcd中对应的service和endpoint信息

在服务器中测试ClusterIP

```

root@aba45a225d81:/home# ip a s
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0@if12: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP group default
    link/ether 02:11:ff:7b:59:67 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.2.16.5/24 brd 10.2.16.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::11:ffff:fe7b:5967/64 scope link
        valid_lft forever preferred_lft forever
root@aba45a225d81:/home# python3 server.py
Server running on port 12345...

```

在node3节点上起了一个container，在12345端口上开启一个服务

```

root@node-1:/home# ipvsadm -A -t 10.10.0.1:8410 -s rr
root@node-1:/home# ip addr add 10.10.0.1/24 dev flannel.1
RTNETLINK answers: File exists
root@node-1:/home# ipvsadm -a -t 10.10.0.1:8410 -r 10.2.16.5:12345 -m
root@node-1:/home# ipvsadm -Ln
IP Virtual Server version 1.2.1 (size=4096)
Prot LocalAddress:Port Scheduler Flags
  -> RemoteAddress:Port          Forward Weight ActiveConn InActConn
TCP  10.10.0.1:8410 rr
    -> 10.2.16.5:12345             Masq    1      0      0
root@node-1:/home# nc -zv 10.10.0.1 8410
Connection to 10.10.0.1 8410 port [tcp/*] succeeded!
root@node-1:/home#

```

```
17: flannel.1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UNKNOWN group default  
link/ether 2a:f1:30:e9:8b:2f brd ff:ff:ff:ff:ff:ff  
inet 10.2.6.0/32 scope global flannel.1  
    valid_lft forever preferred_lft forever  
inet 10.10.0.1/24 scope global flannel.1  
    valid_lft forever preferred_lft forever  
inet6 fe80::28f1:30ff:fee9:8b2f/64 scope link  
    valid_lft forever preferred_lft forever
```

在node1节点上手动执行命令绑定ClusterIP到flannel.1虚拟网卡上，并配置endpoint（真正提供服务的节点）

这里就是配置前面node3的容器中开启的服务

执行nc指令和curl指令发现连接成功，得到node3的容器中服务的返回值Hello, world

```
root@node-1:/home# nc -zv 10.10.0.1 8410  
Connection to 10.10.0.1 8410 port [tcp/*] succeeded!  
root@node-1:/home# curl 10.10.0.1:8410  
Hello, world!root@node-1:/home#
```